

Software Technologies for Data Science

Lecture 19

SQL Constraints and Keys

Ken Cameron

Content

- INSERT with SELECT
- SELECT expressions
- CONSTRAINTS
- KEYS

INSERT with SELECT

- Until now we've either used `INSERT` or `SELECT` in a query.
- It is possible to use both together.
 - This allows us to use data from one or more of the tables in the database
 - To add rows to another.
- Avoids copying the data to the application and then inserting new rows from the application.

INSERT with SELECT

```
INSERT INTO tablex (colx1, colx2, ..., colxn)  
  SELECT coly1, coly2, ..., colyn FROM tabley ...
```

- colx₁, colx₂, ..., colx_n are the columns of tablex.
- coly₁, coly₂, ..., coly_n are the columns produced by the select.

INSERT with SELECT

```
INSERT INTO tablex (colx1, colx2, ..., colxn)  
  SELECT coly1, coly2, ..., coly2 FROM tabley ...
```

- `colx1`, `colx2`, ..., `colxn` are the columns of `tablex`.
- `coly1`, `coly2`, ..., `colyn` are the columns produced by the select.
- Both lists must be the same length.

SELECT result

- Each result item in a SELECT is not restricted to a column name.
- It can be any expression.
 - And it can be as simple as a constant
 - Or as complicated as you need.
 - We've already seen the aggregate functions of GROUP BY.

SELECT expressions

- The SQL language supports all the usual operators.
 - Numeric, string, logical, etc.
- And a long list of useful functions.
 - Look them up.
- You can include any column from a table you have used in the SELECT.

Aggregate Functions

`BIT_AND()`

Bitwise and

`BIT_OR()`

Bitwise or

`BIT_XOR()`

Bitwise xor

`COUNT()`

Number of rows grouped

`COUNT(DISTINCT)`

Number of distinct rows grouped

`GROUP_CONCAT()`

Concatenated strings

`MAX()`

Maximum value

`MIN()`

Minimum value

Aggregate Functions

AVG()

The average value.

STDDEV()

The population standard deviation

STDDEV_POP()

The population standard deviation

STDDEV_SAMP()

The sample standard deviation

VARIANCE()

The population standard variance

VAR_POP()

The population standard variance

VAR_SAMP()

The sample standard variance

Constraints

- SQL allows us to place constraints on tables and columns.
 - This helps ensure the integrity of the data.
- We've already introduced a couple of them:
 - `NOT NULL` is a constraint on the contents of a column.
 - `PRIMARY KEY` is another constraint on the contents of a column.

Other Common Constraints

- `UNIQUE`
 - All entries must be unique.
 - This is also permits `NULL` as a value. But only one instance.
- `INDEX`
 - Indicates we wish to access the table contents via references to this column.
 - And therefore want it to be efficient.

Indexes

- Without an index, if SQL needs to find rows that match a value in a column it might search through the whole table.
- With an index, the search can be optimised.
 - For example, MySQL uses B-trees for most indexes.
 - It may also use R-trees (for spatial data) and hashing.

Other Common Constraints

- CHECK
 - Ensures that values satisfy a condition.
- DEFAULT
 - Used to set a default value for a column when no value is specified.
 - c.f. Setting the parameter in a python function definition.

Example

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    Name varchar(255) NOT NULL,  
    Age int,  
    Member int,  
    City varchar(255),  
    CHECK (Age >= 18),  
    CONSTRAINT Chk_Person CHECK (Member=1 AND  
                                   City='Bath')  
)
```

FOREIGN KEY

- FOREIGN KEY is used to link tables together.
 - A child table contains a FOREIGN KEY.
 - That maps to a PRIMARY KEY in a parent or referenced table.
 - The two keys must be of the same type.
 - The constraint is that a value in the child table is only permitted if it exists in the parent table.

Example

- Consider the `studentid` from our common example.
 - In the `students` table it is the primary key.
- In all the other tables we should define it as a FOREIGN KEY.
 - And map it to the primary key.

Example

MySQL

```
CREATE TABLE enrolled (  
    studentid INT,  
    unitid INT,  
    FOREIGN KEY (studentid) REFERENCES students(studentid),  
    FOREIGN KEY (unitid) REFERENCES units(unitid),  
)
```

SQL Server/Oracle/MS Access

```
CREATE TABLE enrolled (  
    studentid INT FOREIGN KEY REFERENCES students(studentid),  
    unitid INT FOREIGN KEY REFERENCES units(unitid)  
)
```

Other Keys

- PRIMARY, FOREIGN and UNIQUE are not the only types of key.
- There are also
 - Candidate Keys
 - Alternate Keys
 - Super Keys
 - Composite/Compound Keys

Candidate Key

- A Candidate Key is a set of one or more fields/columns that
can identify a record uniquely in a table.
- There can be multiple Candidate Keys in one table.
- Each Candidate Key can work as Primary Key.

Alternate Key

- An alternate key is one that could serve as a PRIMARY KEY.
 - But is n't currently the PRIMARY KEY.
- For example,

In a table that uses the National Insurance Number as the PRIMARY KEY,

The Unique Tax Reference could be suitable Alternate Key.

Composite Key

- Composite Key is a combination of
more than one fields/columns of a table.
- It can be a Candidate Key
- And therefore, can be a Primary key

Super Key

- A Super Key is
 - A set of one or more keys
 - that can be used to identify a record uniquely in a table.
- Primary key, Unique key, Alternate key are subsets of Super Keys.

Why think about keys?

- Keys help us indentify the rows in a table we care about.
- Is every row in the table unique?
 - Do we need it to be?
- How complex does our `WHERE` clause need to be to identify the row(s) we're interested in?
 - Simpler will be faster.

Summary

- INSERT with SELECT
- SELECT expressions
- Constraints
- Keys